



Defense in Depth - The Path to SGX at Akamai

Samuel Erb -- serb@akamai.com
[@erbbysam](#)

DEFCON 26 PACKET HACKING VILLAGE

whoami

Samuel Erb -- @erbbysam

Software Engineer at Akamai

DC23, DC24 black badge (Badge Challenge, Co9)





(disclaimer)

Nothing presented here gives you permission to poke and prod at Akamai.

Everything presented here is the work of **many** individuals within Akamai.

Defense in Depth

Dictionary

defense in depth



defense in depth

phrase of [defense](#)

1. the practice of arranging defensive lines or fortifications so that they can defend each other, especially in case of an enemy incursion.



Translations, word origin, and more definitions

Feedback

Defense in depth (computing) - Wikipedia

[https://en.wikipedia.org/wiki/Defense_in_depth_\(computing\)](https://en.wikipedia.org/wiki/Defense_in_depth_(computing)) ▼

Defense in depth (also known as Castle Approach) is an information assurance (IA) concept in which multiple layers of security controls (**defense**) are placed throughout an information technology (IT) system.

[Background](#) · [Controls](#)

Defense in Depth (My Definition)

1. Avoid the next TLS heartbleed at Akamai
2. Put systems in place to accomplish #1

Heartbleed?



- OpenSSL users were vulnerable to Heartbleed

<https://blogs.akamai.com/2014/04/heartbleed-a-history.html>



RFC 6520

TLS/DTLS Heartbeat Extension

February 2012

When a HeartbeatRequest message is received and sending a HeartbeatResponse is not prohibited as described elsewhere in this document, the receiver **MUST** send a corresponding HeartbeatResponse message carrying an exact copy of the payload of the received HeartbeatRequest.

TLS?

RFC 5246

TLS

August 2008

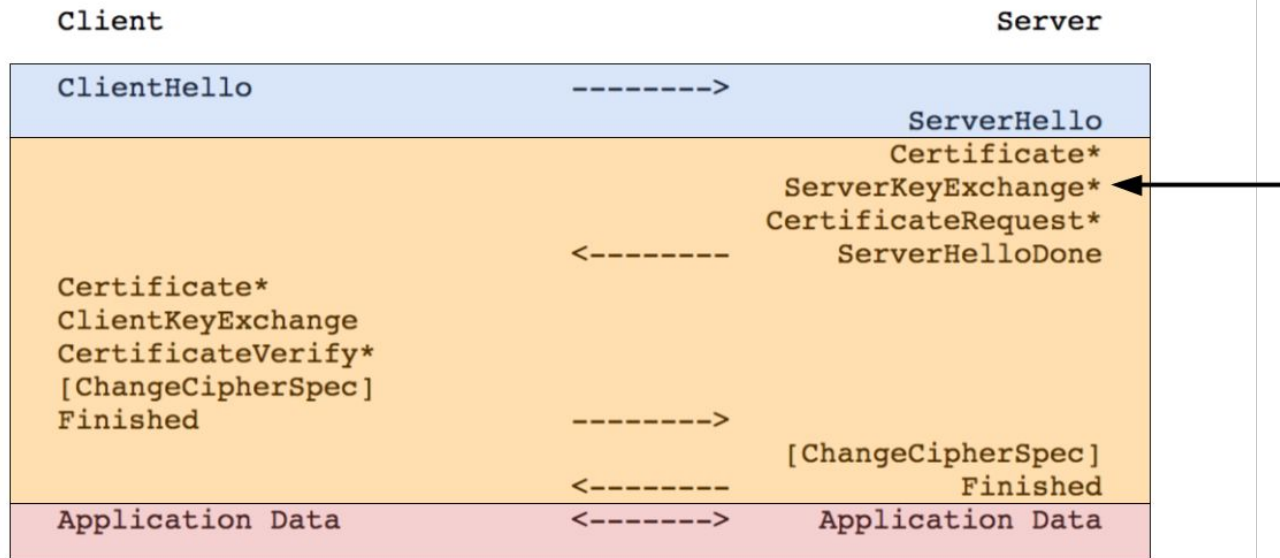


Figure 1. Message flow for a full handshake

* Indicates optional or situation-dependent messages that are not always sent.

TLS Security Journey

- 1) **Make secrets harder to find in memory**

OpenSSL Secure Heap

- One of Akamai's contributions to OpenSSL
- Best effort to protect long lived secrets from overruns, underruns, swapping to disk, appearing in coredumps etc
- Utilizes standard Linux memory protection features

more information - https://www.openssl.org/docs/man1.1.0/crypto/OPENSSL_secure_malloc.html

TLS Security Journey

- 1) Make secrets harder to find in memory
- 2) **Move the secrets out of the process**

OpenSSL Engine Interface & TLS

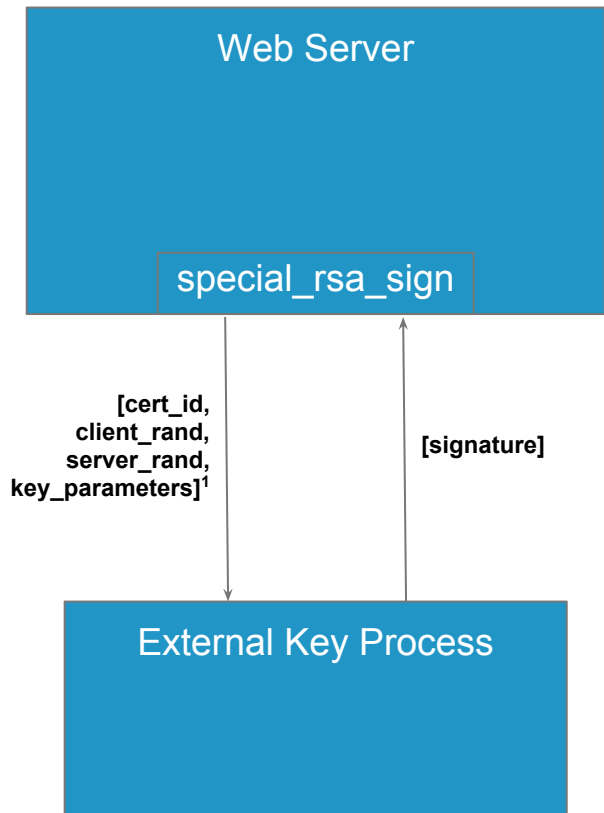
```
static int special_rsa_sign(int flen, const unsigned char *from,
                           unsigned char *to, RSA *rsa, int padding)
{
    int ret;

    // handoff to external RSA processor

    return ret;
}

// this is code for openssl > 1.1, previous versions look slightly different
// you may need this -- #if OPENSSL_VERSION_NUMBER >= 0x10100000L
RSA_METHOD *rsa_method = RSA_meth_new("openssl RSA_METHOD",
                                       RSA_METHOD_FLAG_NO_CHECK);
RSA_meth_set_priv_enc(rsa_method, special_rsa_sign);
ENGINE *e = ENGINE_new();
ENGINE_set_id(e, "rsa-engine");
ENGINE_set_name(e, "rsa engine");
ENGINE_set_RSA(e, rsa_method);
ENGINE_add(e);
ENGINE_set_default_RSA(e);
```

OpenSSL Engine Interface & TLS



```
static int special_rsa_sign(int flen, const unsigned char *from,
                           unsigned char *to, RSA *rsa, int padding)
{
    int ret;

    // handoff to external RSA processor

    return ret;
}
```

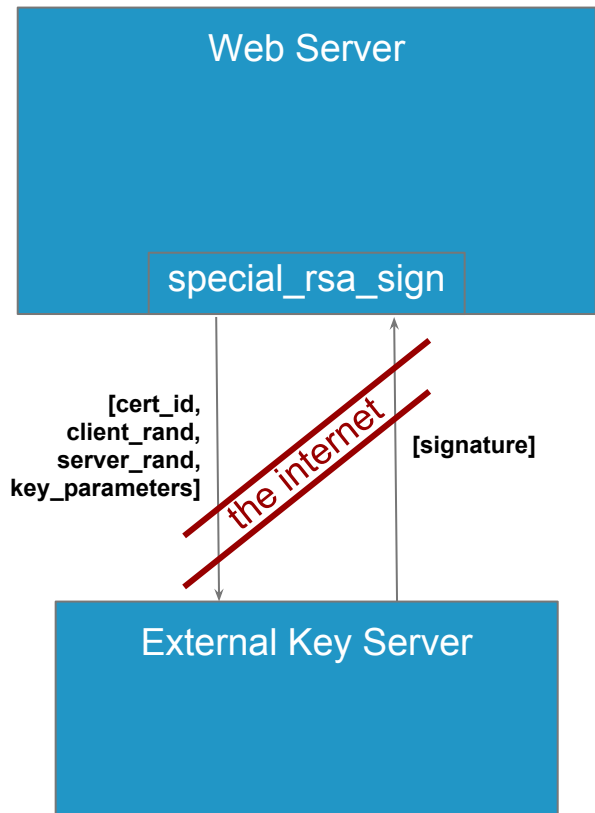
```
// this is code for openssl > 1.1, previous versions look slightly different
// you may need this -- #if OPENSSL_VERSION_NUMBER >= 0x10100000L
RSA_METHOD *rsa_method = RSA_meth_new("openssl RSA METHOD",
                                       RSA_METHOD_FLAG_NO_CHECK);
RSA_meth_set_priv_enc(rsa_method, special_rsa_sign);
ENGINE *e = ENGINE_new();
ENGINE_set_id(e, "rsa-engine");
ENGINE_set_name(e, "rsa engine");
ENGINE_set_RSA(e, rsa_method);
ENGINE_add(e);
ENGINE_set_default_RSA(e);
```

1 - See TLS 1.2 rfc5246 section 7.4.3, "signed_params"

TLS Security Journey

- 1) Make secrets harder to find in memory
- 2) Move the secrets out of the process
- 3) **Move the secrets out of the machine**

Secret Handling Machines



- Protocol is identical to placing secrets in another process
- This creates a many-to-many connection handling & internet routing problem
 - No slow handshakes!
 - There is a performance/security trade off here, at least on the initial TLS handshake

Visualizing a Planetary Scale Network

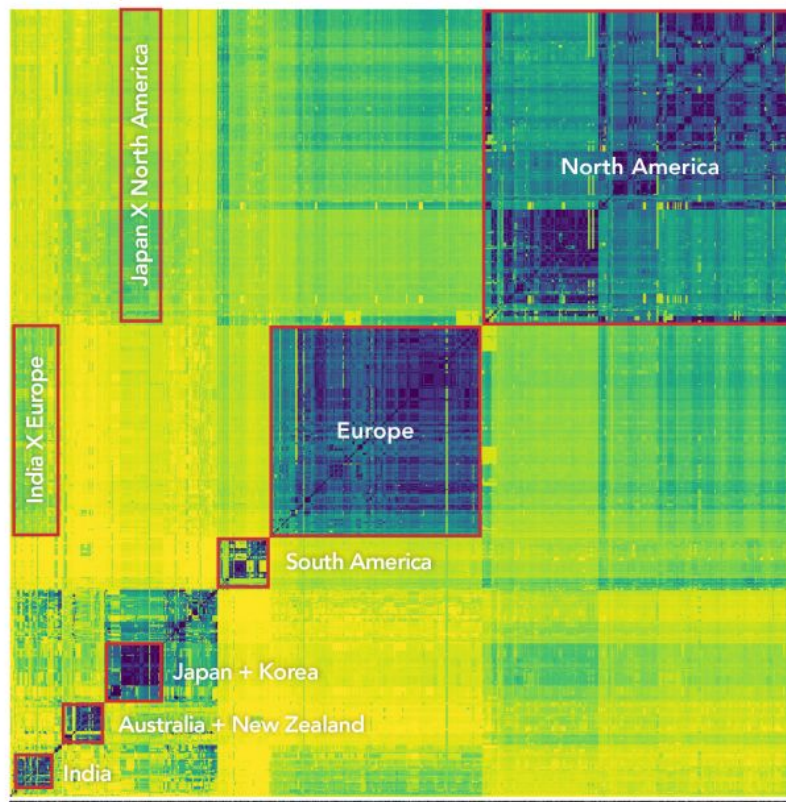
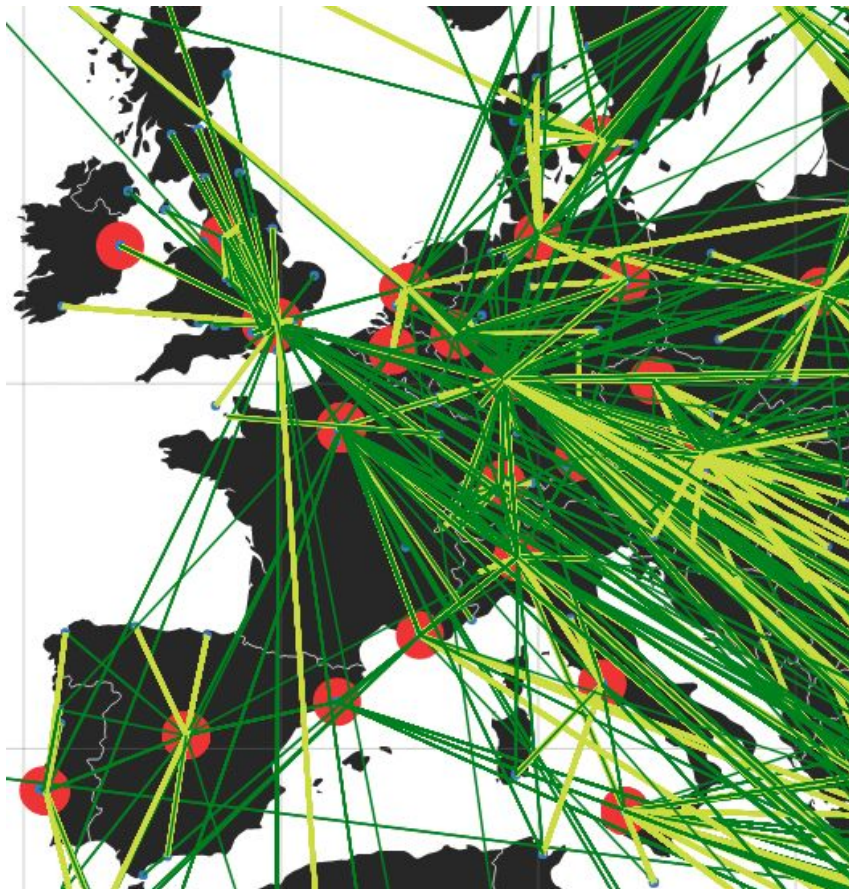


Figure 4-1: Dark colors like purples and blues show short latencies between regions, while light colors (greens and yellows) represent long latencies

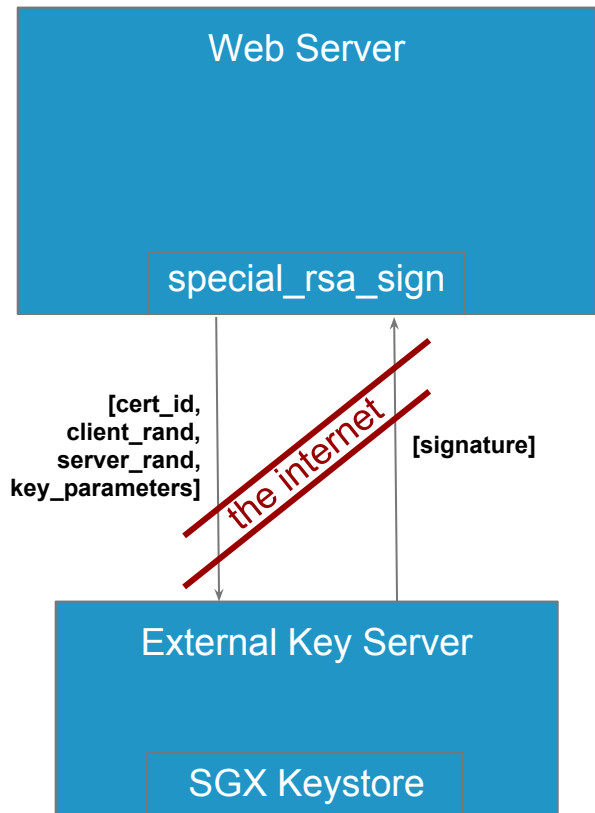
Planetary Scale Network in Practice



TLS Security Journey

- 1) Make secrets harder to find in memory
- 2) Move the secrets out of the process
- 3) Move the secrets out of the machine
- 4) **Move the secrets into an HSM-like device (Intel SGX)**

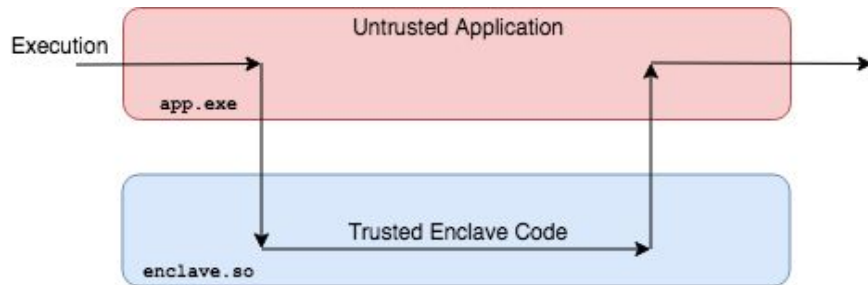
SGX Keystore (HSM-like device)



- Being a bit paranoid -- we've now isolated our secrets, can we remove them from application memory altogether?
- We could also place this SGX keystore inside the web server

What is SGX?

- Software Guard Extensions
- A series of new CPU instructions for handling enclaves
- Enclaves are signed executables that are tightly controlled and cannot be externally monitored
- Entirely on-die - Anything leaving the chip is encrypted
- Extremely limited communication is possible between an enclave and the rest of the system



What is SGX? (continued)

- Reverse sandbox¹
 - Enclaves can read/write to external memory
 - Applications (at any privilege) cannot read/write to enclave memory
- ~Same performance as CPU itself



1 - term originally coined here: <https://www.blackhat.com/docs/us-16/materials/us-16-Aumasson-SGX-Secure-Enclaves-In-Practice-Security-And-Crypto-Review.pdf>
image via <https://commons.wikimedia.org/wiki/File:Sandpit.jpg>

SGX Enclaves (high level)

- Enclaves must be signed. The key used must be trusted by the system.
 - Chain of trust: BIOS trusted “launch enclave” is signed -> your enclave is signed by a key trusted by the “launch enclave” -> “launch enclave” hands out “launch token” allowing your production enclave to start
- Enclaves must be statically compiled
- Requires BIOS & kernel support
- Enclaves cannot interact with anything off-die
- Remote & local attestation is possible



SGX Enclaves (lower level)

- New instructions!
 - Most are for launching and using enclaves
 - 2 new cryptographic operations
 - EREPORT – local attestation generation
Local attestation = (prove an enclave is on the same machine)
 - EGETKEY – returns a key based on input parameters
Allows for “sealing” = encrypting using an enclave specific key and saving outside the enclave for later use (on the same machine)

SGX Sealing

- Allows you to derive a key which is shared between all enclaves that are:
 1. On the same system¹
 2. Signed by the same key²
- Key is persistent (through OS installs etc)



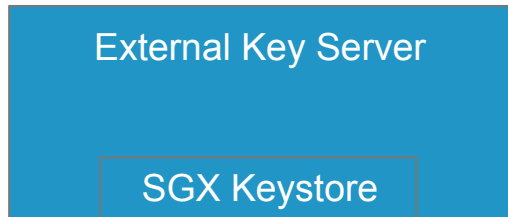
1. (same SEAL_FUSES)

2. (MRSIGNER = signing measurement) -- see See Figure 78 of "Intel SGX explained" <https://eprint.iacr.org/2016/086.pdf>

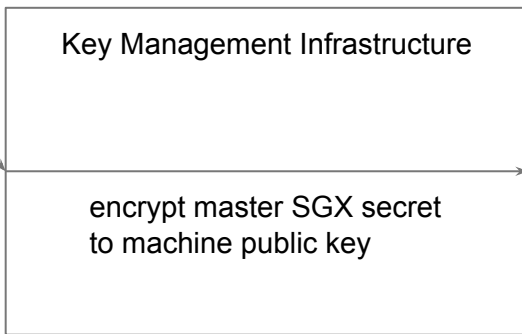
Seal via <http://www.photolib.noaa.gov/htmls/anim0247.htm>

Key Provisioning at Scale

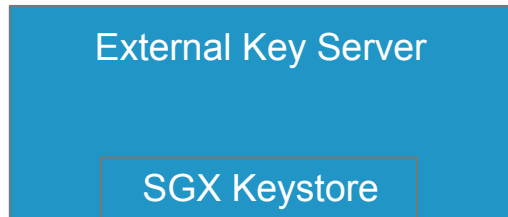
Before Machine Deployment
(trusted environment):



[Sealed private key,
public key]



After Machine Deployment:



[Secrets encrypted to master SGX
secret,
master SGX secret encrypted to
machine public key,
sealed private key, public key]

Complexities with SGX

- Initial trust establishment
- Key rotation, key rotation, key rotation
- Multiple enclaves
- Secure handling of software signing keys
- Integration with SGXSDK & SGXSSL (OpenSSL in an enclave)
- SGX replay attack prevention

Web Server

special_rsa_sign

[cert_id,
client_rand,
server_rand,
key_parameters]

the internet

[signature]

External Key Server

SGX Keystore

Questions?

serb@akamai.com
@erbbysam

